

Assignment1

Name: Yubin Hu

ID: 11712121

1 Part I: the perceptron

Code

```
In [1]: import numpy as np
import matplotlib.pyplot as plt

def gen_gaussian_distribution(size, mean=None, cov=None):
    if not mean:
        mean = np.random.randn(2)
    if not cov:
        cov = np.eye(2)
    data = np.random.multivariate_normal(mean, cov, size)
    return data

class Perceptron(object):

    def __init__(self, n_inputs, max_epochs=1e2, learning_rate=1e-2):
        """
        Initializes perceptron object.
        Args:
            n_inputs: number of inputs.
            max_epochs: maximum number of training cycles.
            learning_rate: magnitude of weight changes at each training cycle
        """
        self.n_inputs = n_inputs
        self.max_epochs = max_epochs
        self.learning_rate = learning_rate
        self.weights = np.zeros(self.n_inputs)
        self.bias = 0

    def forward(self, input):
        """
        Predict label from input
        Args:
            input: array of dimension equal to n_inputs.
        """
        sum = np.sign(np.dot(input, self.weights))
        label = np.where(sum > 0, 1, -1)
        return label

    def train(self, training_inputs, labels):
        """
        Train the perceptron
        Args:
            training_inputs: list of numpy arrays of training points.
            labels: arrays of expected output value for the corresponding points
        """
        train_size = len(training_inputs)
        epochs = 0
        while epochs < self.max_epochs:
            epochs += 1
```

```

        for i in range(train_size):
            if np.any(labels[i] * (np.dot(self.weights, training_inputs[i]
                self.weights = self.weights + (self.learning_rate * labels[i]
                self.bias = self.bias + self.learning_rate * labels[i]

    def score(self, test_inputs, test_labels):
        pred_arr = np.where(self.forward(test_inputs) > 0, 1, -1)
        true_size = len(np.where(pred_arr == test_labels)[0])
        return true_size / len(test_labels)

def main():
    p = Perceptron(2)

    """
    gen dataset
    """
    data_size = 100
    train_size = 80
    x1 = gen_gaussian_distribution(data_size, [5, 5])
    x2 = gen_gaussian_distribution(data_size, [-5, -5])
    y1 = a_label = np.ones(data_size, dtype=np.int16)
    y2 = -y1
    x_train = np.concatenate((x1[:train_size], x2[:train_size]), axis=0)
    y_train = np.concatenate((y1[:train_size], y2[:train_size]), axis=0)
    x_test = np.concatenate((x1[train_size:], x2[train_size:]), axis=0)
    y_test = np.concatenate((y1[train_size:], y2[train_size:]), axis=0)

    """
    train model
    """
    p.train(x_train, y_train)

    """
    test model
    """
    acc = p.score(x_test, y_test)

if __name__ == "__main__":
    main()

```

1.1 Task 1

Generate a dataset of points in R2. To do this, define two Gaussian distributions and sample 100 points from each. Your dataset should then contain a total of 200 points, 100 from each distribution. Keep 80 points per distribution as the training (160 in total), 20 for the test (40 in total).

```

In [2]: """
gen dataset
"""
data_size = 100
train_size = 80
x1 = gen_gaussian_distribution(data_size, [5, 5])
x2 = gen_gaussian_distribution(data_size, [-5, -5])
y1 = a_label = np.ones(data_size, dtype=np.int16)
y2 = -y1
x_train = np.concatenate((x1[:train_size], x2[:train_size]), axis=0)
y_train = np.concatenate((y1[:train_size], y2[:train_size]), axis=0)
x_test = np.concatenate((x1[train_size:], x2[train_size:]), axis=0)

```

```
y_test = np.concatenate((y1[train_size:], y2[train_size:]), axis=0)
```

```
# plt
plt.plot(x1[:,0], x1[:,1], 'x')
plt.plot(x2[:,0], x2[:,1], 'x')
plt.axis('equal')
plt.savefig('./img/fig1.png')
plt.show()
```

```
-----
FileNotFoundError                                Traceback (most recent call last)
<ipython-input-2-6ff8b82d0010> in <module>
    18 plt.plot(x2[:,0], x2[:,1], 'x')
    19 plt.axis('equal')
--> 20 plt.savefig('./img/fig1.png')
    21 plt.show()

~/local/lib/python3.7/site-packages/matplotlib/pyplot.py in savefig(*args, **
kwargs)
    857 def savefig(*args, **kwargs):
    858     fig = gcf()
--> 859     res = fig.savefig(*args, **kwargs)
    860     fig.canvas.draw_idle() # need this if 'transparent=True' to rese
t colors
    861     return res

~/local/lib/python3.7/site-packages/matplotlib/figure.py in savefig(self, fna
me, transparent, **kwargs)
    2309         patch.set_edgecolor('none')
    2310
-> 2311         self.canvas.print_figure(fname, **kwargs)
    2312
    2313         if transparent:

~/local/lib/python3.7/site-packages/matplotlib/backend_bases.py in print_figu
re(self, filename, dpi, facecolor, edgecolor, orientation, format, bbox_inche
s, pad_inches, bbox_extra_artists, backend, **kwargs)
    2215             orientation=orientation,
    2216             bbox_inches_restore=_bbox_inches_restore,
-> 2217             **kwargs)
    2218         finally:
    2219             if bbox_inches and restore_bbox:

~/local/lib/python3.7/site-packages/matplotlib/backend_bases.py in wrapper(*a
rgs, **kwargs)
    1637         kwargs.pop(arg)
    1638
-> 1639         return func(*args, **kwargs)
    1640
    1641     return wrapper


~/local/lib/python3.7/site-packages/matplotlib/backends/backend_agg.py in pri
nt_png(self, filename_or_obj, metadata, pil_kwargs, *args)
    510         mpl.image.imsave(
    511             filename_or_obj, self.buffer_rgba(), format="png", origin=
"upper",
-> 512             dpi=self.figure.dpi, metadata=metadata, pil_kwargs=pil_kwa
rgs)
    513
    514     def print_to_buffer(self):

~/local/lib/python3.7/site-packages/matplotlib/image.py in imsave(fname, arr,
vmin, vmax, cmap, format, origin, dpi, metadata, pil_kwargs)
    1609         pil_kwargs.setdefault("format", format)
    1610         pil_kwargs.setdefault("dpi", (dpi, dpi))
-> 1611         image.save(fname, **pil_kwargs)
    1612
```

1613

```
~/local/lib/python3.7/site-packages/PIL/Image.py in save(self, fp, format, **
params)
    2159         fp = builtins.open(filename, "r+b")
    2160     else:
-> 2161         fp = builtins.open(filename, "w+b")
    2162
    2163     try:
```

FileNotFoundError: [Errno 2] No such file or directory: './img/fig1.png'



We set $\text{mean1} = [5, 5]$ and $\text{mean2} = [-5, -5]$, generate cov1 and cov2 with $\text{np.eve}(2)$ which returns a 2-D array with ones on the diagonal and zeros elsewhere.

1.2 Task 2

Implement the perceptron following the specs in `perceptron.py` and the pseudocode in `perceptronslides.pdf`.

1.3 Task 3

Train the perceptron on the training data (160 points) and test in on the remaining 40 test points. Compute the classification accuracy on the test set.

In [3]:

```
p = Perceptron(2)

"""
gen dataset
"""
data_size = 100
train_size = 80
x1 = gen_gaussian_distribution(data_size, [5, 5])
x2 = gen_gaussian_distribution(data_size, [-5, -5])
y1 = a_label = np.ones(data_size, dtype=np.int16)
y2 = -y1
x_train = np.concatenate((x1[:train_size], x2[:train_size]), axis=0)
y_train = np.concatenate((y1[:train_size], y2[:train_size]), axis=0)
x_test = np.concatenate((x1[train_size:], x2[train_size:]), axis=0)
y_test = np.concatenate((y1[train_size:], y2[train_size:]), axis=0)

"""
train model
"""
p.train(x_train, y_train)
```

```

"""
test model
"""
acc = p.score(x_test, y_test)

print(f'Perceptron test accuracy: {acc * 100}%')

```

Perceptron test accuracy: 100.0%

1.4 Task 4

Experiment with different sets of points (generated as described in Task 1).
What happens during the training if the means of the two Gaussians are too close and/or if their variance is too high?

In [4]:

```

for _ in range(10):
    p = Perceptron(2)

    """
    gen dataset
    """

    data_size = 100
    train_size = 80
    x1 = gen_gaussian_distribution(data_size, [1, 1])
    x2 = gen_gaussian_distribution(data_size, [1, 1])
    y1 = a_label = np.ones(data_size, dtype=np.int16)
    y2 = -y1
    x_train = np.concatenate((x1[:train_size], x2[:train_size]), axis=0)
    y_train = np.concatenate((y1[:train_size], y2[:train_size]), axis=0)
    x_test = np.concatenate((x1[train_size:], x2[train_size:]), axis=0)
    y_test = np.concatenate((y1[train_size:], y2[train_size:]), axis=0)

    """
    train model
    """

    p.train(x_train, y_train)

    """
    test model
    """

    acc = p.score(x_test, y_test)

    print(f'Perceptron test accuracy: {acc * 100}%')

```

```

Perceptron test accuracy: 47.5%
Perceptron test accuracy: 52.5%
Perceptron test accuracy: 50.0%
Perceptron test accuracy: 60.0%
Perceptron test accuracy: 57.49999999999999%
Perceptron test accuracy: 45.0%
Perceptron test accuracy: 55.000000000000001%
Perceptron test accuracy: 47.5%
Perceptron test accuracy: 52.5%
Perceptron test accuracy: 57.49999999999999%

```

We run 10 times for the close Gaussians([1, 1], [1, 1]), and accuracy is lower then 50%.

In [5]:

```

for _ in range(10):
    p = Perceptron(2)

    """
    gen dataset

```

```

"""
data_size = 100
train_size = 80
x1 = gen_gaussian_distribution(data_size, [-5, -5])
x2 = gen_gaussian_distribution(data_size, [1, 1])
y1 = a_label = np.ones(data_size, dtype=np.int16)
y2 = -y1
x_train = np.concatenate((x1[:train_size], x2[:train_size]), axis=0)
y_train = np.concatenate((y1[:train_size], y2[:train_size]), axis=0)
x_test = np.concatenate((x1[train_size:], x2[train_size:]), axis=0)
y_test = np.concatenate((y1[train_size:], y2[train_size:]), axis=0)

"""

train model
"""

p.train(x_train, y_train)

"""

test model
"""

acc = p.score(x_test, y_test)

print(f'Perceptron test accuracy: {acc * 100}%')

```

```

Perceptron test accuracy: 85.0%
Perceptron test accuracy: 97.5%
Perceptron test accuracy: 97.5%
Perceptron test accuracy: 97.5%
Perceptron test accuracy: 97.5%
Perceptron test accuracy: 97.5%
Perceptron test accuracy: 97.5%
Perceptron test accuracy: 100.0%
Perceptron test accuracy: 92.5%
Perceptron test accuracy: 90.0%

```

We run 10 times for the Gaussians with high variance([-5, -5], [1, 1]), and accuracy is higher then the close Gaussians, over 90%.

2 Part II: the mutli-layer perceptron

2.1 Task 1

Implement the MLP architecture by completing the files mlp_numpy.py and modules.py.

```

In [ ]:
from __future__ import absolute_import
from __future__ import division
from __future__ import print_function

from modules import *

class MLP(object):

    def __init__(self, n_inputs, n_hidden, n_classes):
        """
        Initializes multi-layer perceptron object.
        Args:
            n_inputs: number of inputs (i.e., dimension of an input vector).
            n_hidden: list of integers, where each integer is the number of u
            n_classes: number of classes of the classification problem (i.e.,
        """
        self.n_inputs = n_inputs
        self.n_hidden = n_hidden
        self.n_classes = n_classes

```

```

self.layers = []
n_pre = n_inputs
for n_units in n_hidden:
    self.layers.append(Linear(n_pre, n_units))
    self.layers.append(ReLU())
    pre = n_units
self.layers.append(Linear(pre, n_classes))
self.layers.append(SoftMax())

def forward(self, x):
    """
    Predict network output from input by passing it through several layer
    Args:
        x: input to the network
    Returns:
        out: output of the network
    """
    out = x
    for layer in self.layers:
        out = layer.forward(out)
    return out

def backward(self, dout):
    """
    Performs backward propagation pass given the loss gradients.
    Args:
        dout: gradients of the loss
    """
    for layer in self.layers[::-1]:
        dout = layer.backward(dout)
    return dout

if __name__ == '__main__':
    pass

```

```

In [ ]: import numpy as np

class Linear(object):
    def __init__(self, in_features, out_features):
        """
        Module initialisation.
        Args:
            in_features: input dimension
            out_features: output dimension
        TODO:
        1) Initialize weights self.params['weight'] using normal distribution
        std = 0.0001.
        2) Initialize biases self.params['bias'] with 0.
        3) Initialize gradients with zeros.
        """
        mean = 0
        std = 0.0001
        size = (in_features, out_features)
        self.params = {}

        self.params['weight'] = np.random.normal(mean, std, size)
        self.params['bias'] = np.zeros(out_features)
        self.gradients = {}

    def forward(self, x):
        """
        Forward pass (i.e., compute output from input).

```

```

    Args:
        x: input to the module
    Returns:
        out: output of the module
    Hint: Similarly to pytorch, you can store the computed values inside
    and use them in the backward pass computation. This is true for *all*
    """
    self.x = x
    w = self.params['weight']
    b = self.params['bias']
    out = np.dot(x, w) + b
    self.out = out
    return out

def backward(self, dout):
    """
    Backward pass (i.e., compute gradient).
    Args:
        dout: gradients of the previous module
    Returns:
        dx: gradients with respect to the input of the module
    TODO:
    Implement backward pass of the module. Store gradient of the loss with
    layer parameters in self.grads['weight'] and self.grads['bias'].
    """
    # TODO
    self.gradients['weight'] = np.dot(np.transpose(self.x), dout) / self
    self.gradients['bias'] = np.mean(dout, axis=0)
    dx = np.dot(dout, np.transpose(self.params['weight']))
    return dx

class ReLU(object):
    def forward(self, x):
        """
        Forward pass.
        Args:
            x: input to the module
        Returns:
            out: output of the module
        """
        self.x = x;
        out = np.maximum(x, 0)
        return out

    def backward(self, dout):
        """
        Backward pass.
        Args:
            dout: gradients of the previous module
        Returns:
            dx: gradients with respect to the input of the module
        """
        dx = np.where(self.x > 0, dout, 0)
        return dx

class SoftMax(object):
    def exp_normalize(self, x):
        b = x.max()
        y = np.exp(x - b)
        return y / np.reshape(y.sum(axis=1), (-1, 1))

    def forward(self, x):
        """
        Forward pass.
        Args:

```



```

        x: input to the module
Returns:
        out: output of the module

TODO:
Implement forward pass of the module.
To stabilize computation you should use the so-called Max Trick
https://timvieira.github.io/blog/post/2014/02/11/exp-normalize-trick/

"""
self.x = x
out = self.exp_normalize(x)
self.out = out
return out

def backward(self, dout):
    """
    Backward pass.
    Args:
        dout: gradients of the previous module
    Returns:
        dx: gradients with respect to the input of the module
    """
    # TODO
    dx = (dout - np.reshape(np.sum(dout * self.out, 1), [-1, 1])) * self.out
    return dx

class CrossEntropy(object):
    def forward(self, x, y):
        """
        Forward pass.
        Args:
            x: input to the module
            y: labels of the input
        Returns:
            out: cross entropy loss
        """
        # TODO
        out = np.sum(- np.log(np.maximum(x, 1e-8)) * y) / x.shape[0]
        return out

    def backward(self, x, y):
        """
        Backward pass.
        Args:
            x: input to the module
            y: labels of the input
        Returns:
            dx: gradient of the loss with respect to the input x.
        """
        dx = - y / (np.maximum(x, 1e-8))
        return dx

if __name__ == "__main__":
    softmax = SoftMax()

    crossentropy = CrossEntropy()

```

2.2 Task 2

Implement training and testing script in train mlp numpy.py. (Please keep 80% of the dataset for training and the remaining 20% for testing. Note that this is a random split of 80% and

20%)

```
In [3]: def generate_dataset():
        data_size = 1000
        train_size = 800
        x, y = datasets.make_moons(data_size, shuffle=True, noise=None)
        # one hot
        data_shape = 2
        y = np.eye(data_shape)[y.reshape(-1)]
        train_x = x[:train_size]
        train_y = y[:train_size]
        test_x = x[train_size:]
        test_y = y[train_size:]

        return train_x, train_y, test_x, test_y
```

2.3 Task 3

Using the default values of the parameters, report the results of your experiments using a jupyter notebook where you show the accuracy curves for both training and test data.

3 Part III: stochastic gradient descent

3.1 Task 1

Modify the train method in train_mlp_numpy.py to accept a parameter that allows the user to specify if the training has to be performed using batch gradient descent (which you should have implemented in Part II) or stochastic gradient descent.

3.2 Task 2

Using the default values of the parameters, report the results of your experiments using a jupyter notebook where you show the accuracy curves for both training and test data.

How to run code

for SGD, default is BGD

```
python3 train_mlp_numpy.py --gradient_descent SGD
```